

A Public Registry for Intellectual Property in the Age of AI

A protocol for real-time, consent-based use of likeness, voice, and creative work

Concept note · Faheem · July 2026

The problem

When a company trains a large AI model, it draws on essentially the whole public internet including copyrighted images, the faces of public figures, recorded voices, and written work whose owners never agreed to this use. The deeper issue is not only that this data was ingested during training, but that at the moment of generation there is no layer of enforcement between an owner's wishes and the model's output. Even when a rights holder has clearly stated how their image or work may be used or not used, that intent exists in no form the model can query. The gap between what an owner permits and what a model will produce is, today, unbridged.

The idea

A real-time, public registry for intellectual property. A rights holder a public figure, a creator, a company, or an ordinary individual registers their protected assets and attaches a usage policy to each: images and likeness, written work, voice and audio, and other digital property. Before an AI system generates or returns content involving a protected person or work, it consults the registry. If the asset is registered as restricted, the system declines, applies the stated conditions, or provides attribution as required. In effect it is a permissions layer for generative AI the missing bridge between a stated policy and a model's behaviour, checked at the instant it matters.

How it works: decentralised, like the internet or domain names

The registry is designed to operate the way the internet itself operates: as open, distributed infrastructure rather than a product owned by one organisation. The natural reference points are DNS and WHOIS public, interoperable systems of record that anyone can query and that no single party controls. It has three moving parts in sequence: rights holders register their assets and policies; a distributed registry stores these records across many independent nodes synchronised by a common open protocol, with no central point of control; and any AI system queries the registry in real time before producing an output, honouring whatever terms it finds. Because it is an open protocol rather than a private service, it is interoperable across every platform and every size of model. Its resilience comes precisely from not being centralized.

Why build it this way

The primary goal is not commercial gain but give the world a system including community, governance model, technical solution and keep the authorship of something foundational closer to the way the World Wide Web is inseparable from Tim Berners-Lee, who declined to patent it so it could remain a shared resource. That points to a research-led path: document the idea through a university, build a modest proof-of-concept, engage AI-safety groups and rights organisations, govt early, and in time place the work under an independent non-profit or standards body so it stays neutral. Dutch Govt needs to involve to work closely to enforce the generative AI companies to adhere the Intellectual property policy. The idea can implement in EU and in world gradually. The author's role is that of architect and steward.

Open questions I would value guidance on

- Whether the core idea is solid enough to develop further, and where it is weakest.
- The hardest technical problem verifying genuine ownership and preventing false claims.
- How the idea relates to existing frameworks such as the EU AI Act, GDPR, and copyright law.
- The right way to publish and attribute the work so authorship is clearly established.

A fuller draft including possible models for sustaining the registry over time is available if useful, but this note is deliberately limited to the core concept.